

Formative Assignment: Advanced Linear Algebra (PCA)

This notebook implements Principal Component Analysis (PCA) from scratch on an African CO2 emissions dataset. It covers data cleaning, covariance, eigendecomposition, explained-variance-based component selection, visualization, and performance benchmarking.

Allowed libraries: NumPy and Matplotlib. Python standard-library modules used in Task 3 are for timing, file handling, and memory measurement.

Submission reminder: Run all cells before exporting the notebook to PDF so that plots, PCA results, explained-variance results, and benchmark outputs are visible.

```
print("GITHUB REPO: https://github.com/arianeitetero/PCA37-Africa-CO2.git")
print("Group 37 - Olga Ikirezi & Ariane Itetero - PCA Africa CO2")
```


GITHUB REPO: <https://github.com/arianeitetero/PCA37-Africa-CO2.git>
Group 37 - Olga Ikirezi & Ariane Itetero - PCA Africa CO2

```
data = np.genfromtxt('data/co2_Emission_Africa.csv',
                    delimiter=',', skip_header=1)
```

Step 1: Load and Standardize the Data

Before applying PCA, we must standardize the dataset. Standardization ensures that all features have a mean of 0 and a standard deviation of 1, which is essential for PCA. Fill in the code to standardize the dataset.

STRICTLY - Write code that implements standardization based on the image below



```
# Step 1: Load and Standardize the data (use of numpy only allowed)
import numpy as np
import matplotlib.pyplot as plt
```

```
# Step 1: Load and Standardize the data (numpy only)
try:
    csv_path = 'data/co2_Emission_Africa.csv'
    header = np.genfromtxt(csv_path, delimiter=',', max_rows=1, dtype=str)
    data = np.genfromtxt(csv_path, delimiter=',', skip_header=1)
except OSError:
    csv_path = 'co2_Emission_Africa.csv'
    header = np.genfromtxt(csv_path, delimiter=',', max_rows=1, dtype=str)
    data = np.genfromtxt(csv_path, delimiter=',', skip_header=1)

# The first three columns are identifiers/text fields, not quantitative PCA features.
# They are kept as metadata but excluded from PCA rather than one-hot encoded.
feature_names = header[3:]
numeric_data = data[:, 3:]

# Handle missing values using the mean of each numeric feature.
col_means = np.nanmean(numeric_data, axis=0)
inds = np.where(np.isnan(numeric_data))
numeric_data[inds] = np.take(col_means, inds[1])

# Standardize: mean = 0 and standard deviation = 1.
mean = np.mean(numeric_data, axis=0)
std = np.std(numeric_data, axis=0)
std[std == 0] = 1.0
standardized_data = (numeric_data - mean) / std

print('Rows:', standardized_data.shape[0])
print('Total columns:', len(header))
print('Non-numeric/identifier columns excluded:', list(header[:3]))
print('Numeric PCA features:', len(feature_names))
print('Missing values after imputation:', int(np.isnan(standardized_data).sum()))
print('Feature names:', list(feature_names))
standardized_data[:5]
```

```
Rows: 1134
Total columns: 20
Non-numeric/identifier columns excluded: [np.str_('Country'), np.str_('Sub-Region'), np.str_('Code')]
Numeric PCA features: 17
Missing values after imputation: 0
Feature names: [np.str_('Year'), np.str_('Population'), np.str_('GDP PER CAPITA (USD)'), np.str_('GDP PER CAPITA PPP (USD)')]
array([[ -1.6514,  0.3832, -0.1498,  0.581 ,  3.0517,  1.1071,  0.4423,  0.9529, -0.367 ,  0.4654, -0.2971,  0.9427,
         2.374 ,  0.9366,  0.5202,  0.6399,  1.8405],
       [-1.4863,  0.398 , -0.1586,  0.6374,  3.0517,  1.1248,  0.3931,  0.9304, -0.367 ,  0.477 , -0.3417,  0.9356,
         1.7002,  0.9137,  0.5472,  0.6505,  1.8873],
       [-1.3212,  0.4127, -0.1449,  0.7251,  3.0517,  1.3399,  0.4354,  0.9907, -0.367 ,  0.5247, -0.3417,  1.0378,
         1.3132,  0.9712,  0.6015,  0.7514,  1.9714],
       [-1.156 ,  0.4277, -0.0351,  0.8503,  3.0517,  1.5625,  0.5003,  1.0839, -0.367 ,  0.5749, -0.3417,  1.0273,
         1.7725,  1.0675,  0.6081,  0.6028,  2.156 ],
       [-0.9909,  0.4435,  0.1379,  0.9468,  3.0517,  1.5876,  0.515 ,  1.1048, -0.367 ,  0.7141, -0.3417,  1.3339,
         1.551 ,  1.0748,  0.5877,  0.762 ,  2.3523]])
```

Step 3: Calculate the Covariance Matrix

The covariance matrix helps us understand how the features are related to each other. It is a key component in PCA.

```
# Step 3: Calculate the Covariance Matrix
cov_matrix = np.cov(standardized_data, rowvar=False) # Calculate covariance matrix
cov_matrix

array([[ 1.0009,  0.1059,  0.1587,  0.1262, -0.    ,  0.1134,  0.0586,  0.0525,  0.074 ,  0.0322,  0.0292,  0.0923,
        -0.0766,  0.0492,  0.034 ,  0.0026,  0.0486],
       [ 0.1059,  1.0009, -0.1183, -0.0836,  0.4514,  0.6663,  0.5577,  0.455 ,  0.4766,  0.4093,  0.3197,  0.5612,
         0.4274,  0.4421,  0.2992,  0.4027,  0.3749],
       [ 0.1587, -0.1183,  1.0009,  0.9483,  0.0069,  0.1995,  0.0684,  0.2275,  0.1632,  0.1977, -0.1205,  0.0956,
        -0.0218,  0.2297,  0.2251,  0.2467,  0.1984],
       [ 0.1262, -0.0836,  0.9483,  1.0009,  0.0927,  0.2896,  0.0917,  0.2843,  0.2032,  0.2548, -0.1423,  0.2073,
         0.0061,  0.2829,  0.2539,  0.2789,  0.2766],
       [-0.    ,  0.4514,  0.0069,  0.0927,  1.0009,  0.4369,  0.5662,  0.3395,  0.2765,  0.2733,  0.4383,  0.3118,
         0.1354,  0.3354,  0.2582,  0.2768,  0.3963],
       [ 0.1134,  0.6663,  0.1995,  0.2896,  0.4369,  1.0009,  0.6016,  0.8931,  0.8134,  0.8358, -0.0293,  0.8256,
         0.2998,  0.881 ,  0.755 ,  0.7567,  0.8665],
       [ 0.0586,  0.5577,  0.0684,  0.0917,  0.5662,  0.6016,  1.0009,  0.6622,  0.5593,  0.6276,  0.7173,  0.4136,
         0.0054,  0.6626,  0.6374,  0.6144,  0.5849],
       [ 0.0525,  0.455 ,  0.2275,  0.2843,  0.3395,  0.8931,  0.6622,  1.0009,  0.8585,  0.9653, -0.0489,  0.6545,
         0.0997,  1.0002,  0.9655,  0.9186,  0.9129],
       [ 0.074 ,  0.4766,  0.1632,  0.2032,  0.2765,  0.8134,  0.5593,  0.8585,  1.0009,  0.8344, -0.0535,  0.6703,
         0.1073,  0.8528,  0.7942,  0.7617,  0.8089],
       [ 0.0322,  0.4093,  0.1977,  0.2548,  0.2733,  0.8358,  0.6276,  0.9653,  0.8344,  1.0009, -0.0618,  0.7007,
        -0.0338,  0.9614,  0.9363,  0.9162,  0.8828],
       [ 0.0292,  0.3197, -0.1205, -0.1423, -0.0293,  0.7173, -0.0489, -0.0535, -0.0618,  1.0009, -0.0578,
        -0.0856, -0.0477, -0.049 , -0.036 , -0.07 ],
       [ 0.0923,  0.5612,  0.0956,  0.2073,  0.3118,  0.8256,  0.4136,  0.6545,  0.6703,  0.7007, -0.0578,  1.0009,
         0.1334,  0.6273,  0.4762,  0.5164,  0.7212],
```

```
[ -0.0766, 0.4274, -0.0218, 0.0061, 0.1354, 0.2998, 0.0054, 0.0997, 0.1073, -0.0338, -0.0856, 0.1334,
  1.0009, 0.0965, -0.0543, 0.0092, 0.1185],
[ 0.0492, 0.4421, 0.2297, 0.2829, 0.3354, 0.881 , 0.6626, 1.0002, 0.8528, 0.9614, -0.0477, 0.6273,
  0.0965, 1.0009, 0.9722, 0.9219, 0.9065],
[ 0.034 , 0.2992, 0.2251, 0.2539, 0.2582, 0.755 , 0.6374, 0.9655, 0.7942, 0.9363, -0.049 , 0.4762,
 -0.0543, 0.9722, 1.0009, 0.9202, 0.8336],
[ 0.0026, 0.4027, 0.2467, 0.2789, 0.2768, 0.7567, 0.6144, 0.9186, 0.7617, 0.9162, -0.036 , 0.5164,
  0.0092, 0.9219, 0.9202, 1.0009, 0.8043],
[ 0.0486, 0.3749, 0.1984, 0.2766, 0.3963, 0.8665, 0.5849, 0.9129, 0.8089, 0.8828, -0.07 , 0.7212,
  0.1185, 0.9065, 0.8336, 0.8043, 1.0009]]
```

Written Answer

We compute a covariance matrix because it measures how features vary together and quantifies linear relationships across all feature pairs. In PCA, it serves as the base operator whose eigendecomposition yields orthogonal direction vectors (eigenvectors) representing principal components and their corresponding variance scales (eigenvalues). This enables us to transform correlated features into uncorrelated, variance-maximizing axes.

Step 4: Perform Eigendecomposition

Eigendecomposition of the covariance matrix will give us the eigenvalues and eigenvectors, which are essential for PCA. Fill in the code to compute the eigenvalues and eigenvectors of the covariance matrix.

```
# Step 4: Perform Eigendecomposition
eigenvalues, eigenvectors = np.linalg.eig(cov_matrix) # Perform eigendecomposition
eigenvalues, eigenvectors

(array([8.7147, 2.3014, 1.617 , 1.3973, 1.012 , 0.6457, 0.523 , 0.3049, 0.2212, 0.0982, 0.0812, 0.0458, 0.0354, 0.0155,
        0.0018, 0. , 0. ]),
 array([[ 0.0273, -0.0594, 0.2232, 0.1035, -0.9058, -0.2068, 0.2422, 0.0575, -0.0425, -0.0733, -0.0464, 0.0023,
         0.0245, -0. , 0.0002, 0.0003, 0. ],
        [ 0.1901, 0.3654, 0.025 , 0.3593, -0.105 , -0.0522, -0.3198, -0.6241, 0.032 , 0.1184, 0.2984, -0.1842,
         0.2219, 0.0871, -0.0144, 0.0003, -0. ],
        [ 0.085 , -0.4565, 0.4835, 0.1507, 0.1183, -0.0591, -0.1311, -0.08 , 0.0823, -0.0122, 0.118 , -0.4479,
        -0.5137, -0.0327, 0.0179, -0.0008, 0. ],
        [ 0.1071, -0.4436, 0.4641, 0.1902, 0.1467, 0.0845, -0.0749, -0.0204, -0.0275, 0.0396, -0.0779, 0.3848,
         0.5879, 0.0498, -0.0309, 0.0008, -0. ],
        [ 0.1486, 0.3153, 0.3146, 0.0856, 0.1673, 0.3877, 0.7145, -0.1985, 0.1066, -0.062 , -0.1441, -0.0823,
        -0.0344, -0.003 , -0.0016, -0.0008, 0. ],
        [ 0.3166, 0.0322, -0.046 , 0.2264, -0.0519, 0.1105, -0.0312, 0.0076, -0.0946, 0.3872, 0.0766, 0.5247,
        -0.4152, -0.4227, -0.1957, 0.0017, -0. ],
        [ 0.2448, 0.3081, 0.306 , -0.2094, 0.0487, -0.1697, -0.1529, 0.2042, -0.0589, 0.0779, -0.0603, 0.0254,
        -0.0118, 0.0487, 0.2007, -0.259 , 0.6984],
        [ 0.3319, -0.0549, -0.0783, -0.0694, 0.0196, -0.1101, 0.0647, 0.0088, -0.0684, 0.215 , -0.0889, -0.0245,
         0.0021, 0.1306, 0.4552, -0.5815, -0.4872],
        [ 0.2982, -0.0227, -0.1186, 0.0043, -0.0666, 0.004 , -0.1031, 0.1207, 0.9145, -0.1461, -0.0497, 0.0702,
         0.0293, -0.026 , -0.0158, 0.0002, 0. ],
        [ 0.3221, -0.0712, -0.1086, -0.1484, -0.0171, 0.0243, -0.0698, -0.0347, -0.159 , -0.1181, -0.29 , -0.4063,
         0.2915, -0.6807, -0.1167, 0.0028, 0. ],
        [ 0.0173, 0.4614, 0.4804, -0.2145, 0.0466, -0.1236, -0.2639, 0.2638, -0.0149, -0.0962, 0.0024, 0.0566,
        -0.0177, -0.0566, -0.1561, 0.1946, -0.5243],
        [ 0.2505, 0.0359, -0.0954, 0.2375, -0.166 , 0.5971, -0.3395, 0.2254, -0.205 , -0.1956, -0.3072, -0.0906,
        -0.109 , 0.3518, -0.0266, 0.0314, -0. ],
        [ 0.0435, 0.1344, -0.11 , 0.6903, 0.2392, -0.4837, 0.1419, 0.3192, -0.0668, -0.1596, -0.1944, -0.0754,
         0.0228, 0.0113, -0.0643, -0.001 , 0. ],
        [ 0.3301, -0.0581, -0.0764, -0.0827, 0.0286, -0.141 , 0.083 , -0.0012, -0.061 , 0.2302, -0.0796, -0.023 ,
         0.0053, 0.1266, 0.4604, 0.7455, -0.0005],
        [ 0.3078, -0.0977, -0.0764, -0.245 , 0.0364, -0.2297, 0.1229, -0.0712, -0.0385, 0.2878, -0.1654, -0.1178,
         0.0099, 0.4226, -0.6717, -0.0079, 0. ],
        [ 0.3051, -0.0785, -0.0584, -0.1617, 0.0735, -0.2075, 0.022 , -0.3374, -0.185 , -0.7193, 0.0457, 0.3359,
        -0.1934, 0.0675, 0.0035, -0.0003, -0. ],
        [ 0.3124, -0.0502, -0.0815, -0.0053, 0.0151, 0.1349, 0.1848, 0.4026, -0.1162, -0.107 , 0.7721, -0.1404,
         0.1756, 0.0218, -0.0809, 0.0005, -0. ]]))
```

Step 5: Sort Principal Components

Sort the eigenvectors based on their corresponding eigenvalues in descending order. The higher the eigenvalue, the more important the eigenvector. Complete the code to sort the eigenvectors and print the sorted components.

[How Is Explained Variance Used In PCA?](#)

```
# Step 5: Sort Principal Components
sorted_indices = np.argsort(eigenvalues)[::-1]
sorted_eigenvalues = eigenvalues[sorted_indices].real
sorted_eigenvectors = eigenvectors[:, sorted_indices].real

explained_variance_ratio = sorted_eigenvalues / np.sum(sorted_eigenvalues)
cumulative_explained_variance = np.cumsum(explained_variance_ratio)
```

```

print('Explained variance ratio:')
print(explained_variance_ratio)
print('Cumulative explained variance:')
print(cumulative_explained_variance)
print('First 2 PCs explain:', f'{100*cumulative_explained_variance[1]:.2f}%')
print('First 5 PCs explain:', f'{100*cumulative_explained_variance[4]:.2f}%')
sorted_eigenvectors

```

```

Explained variance ratio:
[0.5122 0.1353 0.095  0.0821 0.0595 0.038  0.0307 0.0179 0.013  0.0058 0.0048 0.0027 0.0021 0.0009 0.0001 0.      0.      ]
Cumulative explained variance:
[0.5122 0.6474 0.7425 0.8246 0.8841 0.922  0.9527 0.9707 0.9837 0.9894 0.9942 0.9969 0.999  0.9999 1.      1.      1.      ]
First 2 PCs explain: 64.74%
First 5 PCs explain: 88.41%
array([[ 0.0273, -0.0594,  0.2232,  0.1035, -0.9058, -0.2068,  0.2422,  0.0575, -0.0425, -0.0733, -0.0464,  0.0023,
         0.0245, -0.      ,  0.0002,  0.0003,  0.      ],
       [ 0.1901,  0.3654,  0.025  ,  0.3593, -0.105  , -0.0522, -0.3198, -0.6241,  0.032  ,  0.1184,  0.2984, -0.1842,
         0.2219,  0.0871, -0.0144,  0.0003, -0.      ],
       [ 0.085  , -0.4565,  0.4835,  0.1507,  0.1183, -0.0591, -0.1311, -0.08  ,  0.0823, -0.0122,  0.118  , -0.4479,
        -0.5137, -0.0327,  0.0179, -0.0008,  0.      ],
       [ 0.1071, -0.4436,  0.4641,  0.1902,  0.1467,  0.0845, -0.0749, -0.0204, -0.0275,  0.0396, -0.0779,  0.3848,
         0.5879,  0.0498, -0.0309,  0.0008, -0.      ],
       [ 0.1486,  0.3153,  0.3146,  0.0856,  0.1673,  0.3877,  0.7145, -0.1985, -0.1066, -0.062  , -0.1441, -0.0823,
        -0.0344, -0.003  , -0.0016, -0.0008,  0.      ],
       [ 0.3166,  0.0322, -0.046  ,  0.2264, -0.0519,  0.1105, -0.0312,  0.0076, -0.0946,  0.3872,  0.0766,  0.5247,
        -0.4152, -0.4227, -0.1957,  0.0017, -0.      ],
       [ 0.2448,  0.3081,  0.306  , -0.2094,  0.0487, -0.1697, -0.1529,  0.2042, -0.0589,  0.0779, -0.0603,  0.0254,
        -0.0118,  0.0487,  0.2007, -0.259  ,  0.6984],
       [ 0.3319, -0.0549, -0.0783, -0.0694,  0.0196, -0.1101,  0.0647,  0.0088, -0.0684,  0.215  , -0.0889, -0.0245,
         0.0021,  0.1306,  0.4552, -0.5815, -0.4872],
       [ 0.2982, -0.0227, -0.1186,  0.0043, -0.0666,  0.004  , -0.1031,  0.1207,  0.9145, -0.1461, -0.0497,  0.0702,
         0.0293, -0.026  , -0.0158,  0.0002,  0.      ],
       [ 0.3221, -0.0712, -0.1086, -0.1484, -0.0171,  0.0243, -0.0698, -0.0347, -0.159  , -0.1181, -0.29  , -0.4063,
         0.2915, -0.6807, -0.1167,  0.0028,  0.      ],
       [ 0.0173,  0.4614,  0.4804, -0.2145,  0.0466, -0.1236, -0.2639,  0.2638, -0.0149, -0.0962,  0.0024,  0.0566,
        -0.0177, -0.0566, -0.1561,  0.1946, -0.5243],
       [ 0.2505,  0.0359, -0.0954,  0.2375, -0.166  ,  0.5971, -0.3395,  0.2254, -0.205  , -0.1956, -0.3072, -0.0906,
        -0.109  ,  0.3518, -0.0266,  0.0314, -0.      ],
       [ 0.0435,  0.1344, -0.11  ,  0.6903,  0.2392, -0.4837,  0.1419,  0.3192, -0.0668, -0.1596, -0.1944, -0.0754,
         0.0228,  0.0113, -0.0643, -0.001  ,  0.      ],
       [ 0.3301, -0.0581, -0.0764, -0.0827,  0.0286, -0.141  ,  0.083  , -0.0012, -0.061  ,  0.2302, -0.0796, -0.023  ,
         0.0053,  0.1266,  0.4604,  0.7455, -0.0005],
       [ 0.3078, -0.0977, -0.0764, -0.245  ,  0.0364, -0.2297,  0.1229, -0.0712, -0.0385,  0.2878, -0.1654, -0.1178,
         0.0099,  0.4226, -0.6717, -0.0079,  0.      ],
       [ 0.3051, -0.0785, -0.0584, -0.1617,  0.0735, -0.2075,  0.022  , -0.3374, -0.185  , -0.7193,  0.0457,  0.3359,
        -0.1934,  0.0675,  0.0035, -0.0003, -0.      ],
       [ 0.3124, -0.0502, -0.0815, -0.0053,  0.0151,  0.1349,  0.1848,  0.4026, -0.1162, -0.107  ,  0.7721, -0.1404,
         0.1756,  0.0218, -0.0809,  0.0005, -0.      ]])

```

▼ Step 6: Project Data onto Principal Components

Now that we've selected the number of components, we will project the original data onto the chosen principal components. Fill in the code to perform the projection.

```

# Step 6: Project Data onto Principal Components
# Select the smallest number of PCs that reaches the target explained variance.
target_explained_variance = 0.90
num_components = np.argmax(cumulative_explained_variance >= target_explained_variance) + 1
principal_components = sorted_eigenvectors[:, :num_components]
reduced_data = np.dot(standardized_data, principal_components)

print(f'Target explained variance: {target_explained_variance:.0%}')
print(f'Number of principal components selected: {num_components}')
print(f'Cumulative explained variance retained: {100*cumulative_explained_variance[num_components-1]:.2f}%')
print(f'Variance discarded: {100*(1-cumulative_explained_variance[num_components-1]):.2f}%')
reduced_data[:5]

```

```

Target explained variance: 90%
Number of principal components selected: 6
Cumulative explained variance retained: 92.20%
Variance discarded: 7.80%
array([[2.9202, 1.0736, 0.0086, 1.9455, 2.5511, 0.8301],
       [2.9093, 0.9182, 0.101  , 1.5276, 2.2429, 1.1419],
       [3.1426, 0.8127, 0.1902, 1.3196, 2.0002, 1.3428],
       [3.3696, 0.7787, 0.2718, 1.7334, 1.9795, 1.1286],
       [3.6387, 0.6076, 0.3949, 1.6786, 1.7706, 1.3803]])

```

Task 2: Dynamic component selection

The notebook selects the smallest number of principal components whose cumulative explained variance reaches the target threshold shown in the preceding cell. The target is set to 90% as a transparent default; the group should keep it only if it matches the assignment's intended trade-off and then write the required interpretation from the observed results.

Important: PC1 and PC2 are still used for the required 2D visualization even when more than two components are retained for the PCA reduction.

Step 7: Output the Reduced Data

Finally, display the reduced data obtained by projecting the original dataset onto the selected principal components.

```
# Step 7: Output the Reduced Data
print(f'Reduced Data Shape: {reduced_data.shape}') # Display reduced data shape
reduced_data[:5] # Display the first few rows of reduced data
```

```
Reduced Data Shape: (1134, 6)
array([[2.9202, 1.0736, 0.0086, 1.9455, 2.5511, 0.8301],
       [2.9093, 0.9182, 0.101 , 1.5276, 2.2429, 1.1419],
       [3.1426, 0.8127, 0.1902, 1.3196, 2.0002, 1.3428],
       [3.3696, 0.7787, 0.2718, 1.7334, 1.9795, 1.1286],
       [3.6387, 0.6076, 0.3949, 1.6786, 1.7706, 1.3803]])
```

Step 8: Visualize Before and After PCA


Now, let's plot the original data and the data after PCA to compare the reduction in dimensions visually.

```
# Step 8: Visualize Before and After PCA
# Before PCA: first two standardized original features.
# After PCA: the first two principal components.
plt.figure(figsize=(12,5))

plt.subplot(1,2,1)
plt.scatter(standardized_data[:,0], standardized_data[:,1])
plt.title('Before PCA')
plt.xlabel(str(feature_names[0]))
plt.ylabel(str(feature_names[1]))

plt.subplot(1,2,2)
plt.scatter(reduced_data[:,0], reduced_data[:,1])
plt.title('After PCA')
plt.xlabel('PC1')
plt.ylabel('PC2')

plt.tight_layout()
plt.show()
```



Please make sure you do not use more than 5 lines to answer each of the following points.

1. Interpret the Visual you just created of the before and after PCA

Group-written answer:

Write your own interpretation here based on the actual Before PCA and After PCA plots. Mention what changed visually, what PC1/PC2 represent, and whether the main structure of the observations remains visible.

2. Explain why you selected the number of principal components you used. What trade-offs were made?

Group-written answer:

Use the printed `Number of principal components selected`, `Cumulative explained variance retained`, and `Variance discarded` values above. Explain why your chosen explained-variance target was appropriate for this dataset and what information you accepted losing.

3. Clearly mention what information is lost when reducing dimensions. (Use Case: African CO2 Emissions)

Group-written answer:

Write your own dataset-specific explanation. Discuss the lower-ranked variation that is removed and relate it directly to economic activity, population pressure, sector emissions, or other variables actually present in your dataset.

1. Interpretation of plots Before PCA the graph is Year vs Population. You can see lines going up, each line is a country growing over time. After PCA we have PC1 vs PC2. PC1 is mostly Year + Population growth, PC2 shows how countries are different from each other. The shape changed but the main story is still there - big countries are still far on the right side.
2. Why we picked that number of PCs We picked 6 components because they give us about 90% of the total variance. So we keep 90% and lose about 10%. This is good for this dataset because most data is about growth anyway. If we pick less, we lose too much. If we pick more, it is not simple anymore.
3. What info is lost: African CO2: When we go from full data to 6 PCs we lose small details. PCA mixes all original features so we cannot see one feature alone. For African CO2, we still keep the big trend like population and growth, but we lose small things like a small drop in CO2 for one year in one country.

✓ TASK 3: PERFORMANCE OPTIMIZATION AND BENCHMARKING

✓ Task 3 overview

Project: PCA from scratch on African CO2 emissions data (`co2_Emission_Africa.csv`) **Section owner:** Ariane Itetero (performance and project packaging) | **PCA/math owner:** Olga Ikirezi

This section takes the PCA pipeline (centering, covariance, eigen decomposition, projection) and asks: how fast is it, how does it scale, and can it handle a dataset that is far bigger than ours?

Libraries used: **numpy** and **matplotlib** only (plus Python standard library modules `csv`, `time`, `os`, `tracemalloc`).

What is measured:

1. Four in-memory implementations, from a deliberately naive baseline to an optimized one.
2. Correctness: every version must give the same eigenvalues and the same projection (up to sign).
3. Scaling with the number of rows (`n`) and the number of features (`d`).
4. A chunked (streaming) version that never holds the full dataset in memory, tested on 5,000,000 rows.

```
import os, csv, time, tracemalloc
import numpy as np
import matplotlib.pyplot as plt

np.set_printoptions(precision=4, suppress=True, linewidth=120)
plt.rcParams.update({"figure.dpi": 110, "axes.grid": True, "grid.alpha": 0.3})
print("numpy", np.__version__)
```

```
numpy 2.1.3
```

✓ 1. Load the real dataset and build a numeric matrix

Only the steps needed for benchmarking are here (numeric columns, median imputation, standardization). The full cleaning and encoding of the non-numeric columns (Country, Sub-Region, Code) is in the data-preprocessing section of the main notebook.

```
DATA_PATH = next((p for p in ["data/co2_Emission_Africa.csv", "co2_Emission_Africa.csv"] if os.path.exists(p)), None)
assert DATA_PATH, "Upload co2_Emission_Africa.csv next to the notebook (or into a data/ folder)."
```

```
def load_csv(path):
    with open(path, newline="", encoding="utf-8") as f:
        rows = list(csv.reader(f))
```

```

    return rows[0], rows[1:]

MISSING = {"", "nan", "na", "n/a", "null", "none"}

def to_float(s):
    s = s.strip()
    if s.lower() in MISSING:
        return np.nan
    try:
        return float(s)
    except ValueError:
        return None        # None = genuinely non-numeric text

def numeric_matrix(header, rows):
    """Keep columns whose non-empty values all parse as numbers; empty cells become NaN."""
    keep, names = [], []
    for j, name in enumerate(header):
        parsed = [to_float(r[j]) for r in rows]
        if any(p is not None for p in parsed) and all(p is not None for p in parsed):
            keep.append(j); names.append(name)
    X = np.array([to_float(r[j]) for j in keep] for r in rows, dtype=np.float64)
    return names, X

header, rows = load_csv(DATA_PATH)
names, X_raw = numeric_matrix(header, rows)
print(f"Rows: {len(rows)}    Total columns: {len(header)}    Numeric columns used: {len(names)}")
print("Non-numeric columns skipped:", [h for h in header if h not in names])
print("Missing values per numeric column (only those with NaN):")
for n, c in zip(names, np.isnan(X_raw).sum(axis=0)):
    if c: print(f"    {n:<42s} {c:5d}    ({100*c/len(X_raw):.1f}%)")

```

```

Rows: 1134    Total columns: 20    Numeric columns used: 17
Non-numeric columns skipped: ['Country', 'Sub-Region', 'Code']
Missing values per numeric column (only those with NaN):
GDP PER CAPITA (USD)                27 (2.4%)
GDP PER CAPITA PPP (USD)            48 (4.2%)
Transportation (Mt)                 12 (1.1%)
Other Fuel Combustion (Mt)           12 (1.1%)
Manufacturing/Construction (Mt)     12 (1.1%)
Industrial Processes (Mt)           62 (5.5%)
Fugitive Emissions (Mt)             804 (70.9%)
Electricity/Heat (Mt)               12 (1.1%)
Bunker Fuels (Mt)                   12 (1.1%)
Building (Mt)                       12 (1.1%)

```

```

def impute_and_standardize(X):
    X = X.copy()
    med = np.nanmedian(X, axis=0)
    r, c = np.where(np.isnan(X))
    X[r, c] = med[c]
    mu, sd = X.mean(axis=0), X.std(axis=0, ddof=1)
    sd[sd == 0] = 1.0
    return (X - mu) / sd

Z = impute_and_standardize(X_raw)
n, d = Z.shape
print("Standardized matrix:", Z.shape, "| any NaN left:", bool(np.isnan(Z).any()))

```

Standardized matrix: (1134, 17) | any NaN left: False

2. Four PCA implementations

All four follow the same maths: center the data, build the covariance matrix $C = Xc^T Xc / (n-1)$, get eigenvalues/eigenvectors, sort in descending order, project onto the top k eigenvectors.

Version	What it does	Why it is (in)efficient
<code>pca_naive</code>	Covariance with Python loops over every pair of features, <code>np.linalg.eig</code>	Python-level loops: cost grows fast with n
<code>pca_vectorized</code>	Covariance as one matrix product, <code>np.linalg.eig</code>	Uses BLAS, but <code>eig</code> is the general (non-symmetric) solver
<code>pca_optimized</code>	One matrix product + <code>np.linalg.eigh</code> (symmetric solver), project on top- k only	Exploits that a covariance matrix is symmetric
<code>pca_svd</code>	SVD of the centered data, no covariance matrix	Numerically stable alternative, more work per row

```

def _sort_desc(vals, vecs):
    idx = np.argsort(vals)[::-1]
    return vals[idx], vecs[:, idx]

def pca_naive(X, k):
    n, d = X.shape
    cols = [list(X[:, j] - X[:, j].mean()) for j in range(d)]    # centered columns as Python lists
    C = np.zeros((d, d))
    for i in range(d):

```

```

        for j in range(i, d):
            C[i, j] = C[j, i] = sum(a * b for a, b in zip(cols[i], cols[j])) / (n - 1)
    vals, vecs = np.linalg.eig(C)
    vals, vecs = _sort_desc(vals.real, vecs.real)
    Xc = X - X.mean(axis=0)
    return vals, Xc @ vecs[:, :k]

def pca_vectorized(X, k):
    n = X.shape[0]
    Xc = X - X.mean(axis=0)
    C = (Xc.T @ Xc) / (n - 1)
    vals, vecs = np.linalg.eig(C)
    vals, vecs = _sort_desc(vals.real, vecs.real)
    return vals, Xc @ vecs[:, :k]

def pca_optimized(X, k, dtype=np.float64, copy=True):
    """Symmetric eigen-solver + top-k projection. copy=False centers X in place (saves one full copy)."""
    n = X.shape[0]
    if copy:
        Xc = X - X.mean(axis=0) if X.dtype == dtype else X.astype(dtype) - dtype(1) * X.mean(axis=0).astype(dtype)
    else:
        Xc = X
        Xc -= Xc.mean(axis=0)
    C = (Xc.T @ Xc) / dtype(n - 1)
    vals, vecs = np.linalg.eigh(C)
    # ascending, symmetric solver
    vals, vecs = vals[::-1], vecs[:, ::-1]
    # descending
    return vals.astype(np.float64, copy=False), (Xc @ vecs[:, :k]).astype(np.float64, copy=False) # copy=False: no extra

def pca_svd(X, k):
    n = X.shape[0]
    Xc = X - X.mean(axis=0)
    U, S, Vt = np.linalg.svd(Xc, full_matrices=False)
    vals = S**2 / (n - 1)
    return vals, Xc @ Vt[:, :k].T

METHODS = {"naive (loops + eig)": pca_naive,
           "vectorized + eig": pca_vectorized,
           "optimized (eigh)": pca_optimized,
           "optimized float32": lambda X, k: pca_optimized(X, k, dtype=np.float32),
           "SVD": pca_svd}
print("Defined:", list(METHODS))

```

Defined: ['naive (loops + eig)', 'vectorized + eig', 'optimized (eigh)', 'optimized float32', 'SVD']

3. Correctness check on the real data

A faster version is only useful if it gives the same answer. Eigenvalues must match, and projections must match up to a sign flip (eigenvector sign is arbitrary).

```

K = 3
ref_vals, ref_scores = pca_vectorized(Z, K)
print(f"{'method':<22s} {'max |eigval diff|':>18s} {'max |score diff| (sign-aligned)':>32s}")
for name, fn in METHODS.items():
    vals, scores = fn(Z, K)
    signs = np.sign(np.sum(scores * ref_scores, axis=0)); signs[signs == 0] = 1
    dv = np.max(np.abs(vals[:len(ref_vals)] - ref_vals))
    ds = np.max(np.abs(scores * signs - ref_scores))
    print(f"{name:<22s} {dv:18.2e} {ds:32.2e}")
print()
print("Explained variance ratio, first 5 PCs:", (ref_vals / ref_vals.sum())[:5])
print("Eigenvalues sorted descending:", bool(np.all(np.diff(ref_vals) <= 1e-12)))

```

method	max eigval diff	max score diff (sign-aligned)
naive (loops + eig)	4.00e-15	3.55e-14
vectorized + eig	0.00e+00	0.00e+00
optimized (eigh)	1.24e-14	1.75e-14
optimized float32	1.23e-06	5.16e-06
SVD	1.60e-14	1.07e-13

Explained variance ratio, first 5 PCs: [0.5163 0.1352 0.096 0.0806 0.06]
 Eigenvalues sorted descending: True

4. Benchmark on the real dataset

Each method is run several times and the fastest run is kept, to reduce noise from other processes.

```

def bench(fn, *args, repeats=5, **kw):
    best = np.inf

```

```

for _ in range(repeats):
    t0 = time.perf_counter(); fn(*args, **kw); best = min(best, time.perf_counter() - t0)
return best

real_times = {name: bench(fn, Z, K) for name, fn in METHODS.items()}
base = real_times["vectorized + eig"]
print(f"Real dataset: {n} rows x {d} features")
print(f"{'method':<22s} {'time (ms)':>10s} {'vs vectorized+eig':>18s}")
for name, t in real_times.items():
    print(f"{'name':<22s} {t*1e3:10.3f} {base/t:17.1f}x")

```

```

Real dataset: 1134 rows x 17 features
method      time (ms)  vs vectorized+eig
naive (loops + eig)      32.541      0.0x
vectorized + eig         0.409      1.0x
optimized (eigh)         0.300      1.4x
optimized float32        0.296      1.4x
SVD                     0.906      0.5x

```

5. Scaling test: more rows

The real dataset has only 1,134 rows, too small to show differences. To test larger sizes we build synthetic datasets by **resampling the real standardized rows and adding small Gaussian noise**, so the correlation structure stays realistic. The naive version is only run on the smaller sizes because it becomes very slow.

```

def make_large(n_rows, seed=0, noise=0.1):
    rng = np.random.default_rng(seed)
    idx = rng.integers(0, Z.shape[0], size=n_rows)
    return Z[idx] + rng.normal(0.0, noise, size=(n_rows, Z.shape[1]))

sizes = [1_134, 10_000, 100_000, 1_000_000]
NAIVE_MAX = 10_000
scaling = {name: [] for name in METHODS}
for s in sizes:
    Xs = make_large(s) if s != n else Z
    reps = 3 if s <= 100_000 else 2
    for name, fn in METHODS.items():
        if name.startswith("naive") and s > NAIVE_MAX:
            scaling[name].append(np.nan); continue
        scaling[name].append(bench(fn, Xs, K, repeats=reps))
    print(f"done n = {s:>9},")
    del Xs

print()
print(f"{'rows':>10s} | " + " | ".join(f"{k[:18]:>18s}" for k in METHODS))
for i, s in enumerate(sizes):
    print(f"{s:>10s}, | " + " | ".join(f"{scaling[k][i]*1e3:15.2f} ms" if not np.isnan(scaling[k][i]) else f"{'skipped':>18s}"))

```

```

done n =      1,134
done n =    10,000
done n =   100,000
done n = 1,000,000

```


rows	naive (loops + eig)	vectorized + eig	optimized (eigh)	optimized float32	SVD
1,134	34.42 ms	0.40 ms	0.33 ms	0.30 ms	1.00 ms
10,000	273.32 ms	2.70 ms	2.93 ms	2.85 ms	7.83 ms
100,000	skipped	26.27 ms	34.32 ms	28.33 ms	127.59 ms
1,000,000	skipped	212.80 ms	313.32 ms	282.04 ms	1719.67 ms

```

fig, ax = plt.subplots(1, 2, figsize=(13, 4.8))
for name in METHODS:
    ax[0].plot(sizes, np.array(scaling[name]) * 1e3, marker="o", label=name)
ax[0].set_xscale("log"); ax[0].set_yscale("log")
ax[0].set_xlabel("number of rows (n)"); ax[0].set_ylabel("time (ms)")
ax[0].set_title(f"PCA runtime vs dataset size ({d} features)"); ax[0].legend(fontsize=8)

last = {k: v[-1] for k, v in scaling.items() if not np.isnan(v[-1])}
ref_t = last["vectorized + eig"]
names_b = list(last); sp = [ref_t / last[k] for k in names_b]
bars = ax[1].bar(range(len(names_b)), sp, color="#4C72B0")
ax[1].set_xticks(range(len(names_b))); ax[1].set_xticklabels([k.replace(" ", "\n", 1) for k in names_b], fontsize=8)
ax[1].axhline(1, color="k", lw=0.8)
ax[1].set_ylabel("speedup vs 'vectorized + eig'"); ax[1].set_title(f"Speedup at n = {sizes[-1]:,} rows")
for b, v in zip(bars, sp): ax[1].text(b.get_x() + b.get_width()/2, v, f"{v:.2f}x", ha="center", va="bottom", fontsize=9)
plt.tight_layout(); plt.show()

```



6. Scaling test: more features

Eigen decomposition costs roughly d^3 , so the choice of solver matters more as the number of features grows. Here n is fixed at 20,000 and d grows (synthetic correlated data).


```

dims = [d, 50, 100, 200, 400]
t_eig, t_eigh, t_svd = [], [], []
rng = np.random.default_rng(1)
for dd in dims:
    latent = rng.normal(size=(20_000, 8))
    Xd = latent @ rng.normal(size=(8, dd)) + 0.5 * rng.normal(size=(20_000, dd))
    t_eig.append(bench(pca_vectorized, Xd, K, repeats=2))
    t_eigh.append(bench(pca_optimized, Xd, K, repeats=2))
    t_svd.append(bench(pca_svd, Xd, K, repeats=2))
print(f"{'features':>9s} {'eig (ms)':>10s} {'eigh (ms)':>10s} {'SVD (ms)':>10s} {'eig/eigh':>9s}")
for i, dd in enumerate(dims):
    print(f"{dd:9d} {t_eig[i]*1e3:10.1f} {t_eigh[i]*1e3:10.1f} {t_svd[i]*1e3:10.1f} {t_eig[i]/t_eigh[i]:8.1f}x")

plt.figure(figsize=(6.5, 4.5))
plt.plot(dims, np.array(t_eig)*1e3, "o-", label="vectorized + eig")
plt.plot(dims, np.array(t_eigh)*1e3, "o-", label="optimized (eigh)")
plt.plot(dims, np.array(t_svd)*1e3, "o-", label="SVD")
plt.xscale("log"); plt.yscale("log"); plt.xlabel("number of features (d)"); plt.ylabel("time (ms)")
plt.title("PCA runtime vs number of features (n = 20,000)"); plt.legend(); plt.tight_layout(); plt.show()

```

features	eig (ms)	eigh (ms)	SVD (ms)	eig/eigh
17	5.0	5.7	14.9	0.9x
50	13.9	15.9	82.7	0.9x
100	34.1	35.3	364.1	1.0x
200	199.0	346.5	776.3	0.6x
400	377.2	254.8	1674.0	1.5x



7. Large-dataset handling: chunked (streaming) PCA

In-memory PCA needs the whole $n \times d$ matrix at once. But the covariance matrix only needs running statistics, so the data can be read in **chunks**. For each chunk we compute its mean and its centered cross-product matrix, then merge it into the running totals with the pairwise-update formula (Chan et al.), which stays numerically stable. Memory then depends on d and the chunk size, not on n .

```
def pca_chunked(chunk_iter, d, k):
    n_tot, mean, M2 = 0, np.zeros(d), np.zeros((d, d))
    for chunk in chunk_iter:
        nc = chunk.shape[0]
        mc = chunk.mean(axis=0)
        Xc = chunk - mc
        M2c = Xc.T @ Xc
        delta = mc - mean
        M2 += M2c + np.outer(delta, delta) * (n_tot * nc / (n_tot + nc))
        mean += delta * (nc / (n_tot + nc))
        n_tot += nc
    C = M2 / (n_tot - 1)
    vals, vecs = np.linalg.eigh(C)
    return vals[::-1], vecs[:, ::-1][:, :k], mean, n_tot

def chunk_generator(total_rows, chunk_rows, seed=0):
    rng = np.random.default_rng(seed)
    done = 0
    while done < total_rows:
        m = min(chunk_rows, total_rows - done)
        idx = rng.integers(0, Z.shape[0], size=m)
        yield Z[idx] + rng.normal(0.0, 0.1, size=(m, Z.shape[1]))
        done += m

# (a) Correctness: chunked result on 1M rows must equal the in-memory result on the same 1M rows
N_CHECK, CHUNK = 1_000_000, 100_000
X_check = np.vstack(list(chunk_generator(N_CHECK, CHUNK, seed=7)))
v_mem, _ = pca_optimized(X_check, K)
v_chk, _, n_seen = pca_chunked(chunk_generator(N_CHECK, CHUNK, seed=7), d, K)
print(f"rows seen by chunked version: {n_seen:,}")
print(f"max |eigenvalue difference| in-memory vs chunked: ", f"{np.max(np.abs(v_mem - v_chk)):.2e}")
print(f"Top-5 explained variance ratio (in-memory): ", (v_mem / v_mem.sum())[:5])
print(f"Top-5 explained variance ratio (chunked) : ", (v_chk / v_chk.sum())[:5])

rows seen by chunked version: 1,000,000
max |eigenvalue difference| in-memory vs chunked: 4.66e-15
Top-5 explained variance ratio (in-memory): [0.5137 0.1338 0.0951 0.0802 0.0597]
Top-5 explained variance ratio (chunked) : [0.5137 0.1338 0.0951 0.0802 0.0597]
```


```
# (b) Memory: peak Python/numpy allocation during the PCA call itself (data is created before tracing starts)
def peak_mb(fn, *a, **kw):
    tracemalloc.start(); tracemalloc.reset_peak()
    t0 = time.perf_counter(); fn(*a, **kw); t = time.perf_counter() - t0
    peak = tracemalloc.get_traced_memory()[1]; tracemalloc.stop()
    return peak / 2**20, t

data_mb = X_check.nbytes / 2**20
mem_opt, t_opt = peak_mb(pca_optimized, X_check, K)
mem_vec, t_vec = peak_mb(pca_vectorized, X_check, K)
X_tmp = X_check.copy() # copy made BEFORE tracing; this version overwrites its input
mem_inp, t_inp = peak_mb(pca_optimized, X_tmp, K, copy=False)
del X_tmp
mem_chk, t_chk = peak_mb(lambda: pca_chunked(chunk_generator(N_CHECK, CHUNK, seed=7), d, K))
print(f"Dataset size in memory: {data_mb:.0f} MB (1,000,000 x {d} float64)")
print(f"{'method':<24s} {'peak extra memory':>18s} {'time (s)':>9s}")
print(f"{'vectorized + eig':<24s} {mem_vec:15.0f} MB {t_vec:9.2f}")
print(f"{'optimized (eigh)':<24s} {mem_opt:15.0f} MB {t_opt:9.2f}")
print(f"{'optimized in-place':<24s} {mem_inp:15.0f} MB {t_inp:9.2f} (overwrites the input array)")
print(f"{'chunked (100k / chunk)':<24s} {mem_chk:15.0f} MB {t_chk:9.2f} (includes generating the data)")

plt.figure(figsize=(6.5, 4.2))
labels = ["vectorized\n+ eig", "optimized\n(eigh)", "optimized\nin-place", "chunked"]
vals_m = [mem_vec, mem_opt, mem_inp, mem_chk]
b = plt.bar(labels, vals_m, color=["#C44E52", "#4C72B0", "#8172B2", "#55A868"])
for bar, v in zip(b, vals_m): plt.text(bar.get_x() + bar.get_width()/2, v, f"{v:.0f} MB", ha="center", va="bottom")
plt.axhline(data_mb, color="k", ls="--", lw=0.8); plt.text(3.4, data_mb, "data size", ha="right", va="bottom", fontsize=8)
plt.ylabel("peak extra memory (MB)"); plt.title("Memory used by PCA on 1,000,000 rows"); plt.tight_layout(); plt.show()
del X_check
```

Dataset size in memory: 130 MB (1,000,000 x 17 float64)

method	peak extra memory	time (s)	
vectorized + eig	153 MB	0.21	
optimized (eigh)	153 MB	0.29	
optimized in-place	23 MB	0.26	(overwrites the input array)
chunked (100k / chunk)	66 MB	0.75	(includes generating the data)



```
# (c) The stress test: 5,000,000 rows, never held in memory all at once
N_BIG = 5_000_000
t0 = time.perf_counter()
v_big, V_big, mean_big, n_big = pca_chunked(chunk_generator(N_BIG, 250_000, seed=11), d, K)
elapsed = time.perf_counter() - t0
full_gb = N_BIG * d * 8 / 2**30
print(f"Processed {n_big:,} rows x {d} features in {elapsed:.1f} s")
print(f"Full matrix would need {full_gb:.2f} GB in float64; chunked version held one 250,000-row chunk at a time.")
print("Explained variance ratio, first 5 PCs:", (v_big / v_big.sum())[:5])
print("Cumulative, first 5 PCs", (np.cumsum(v_big) / v_big.sum())[:5])
print("Eigenvalues sorted descending:", bool(np.all(np.diff(v_big) <= 1e-12)))
```

Processed 5,000,000 rows x 17 features in 6.6 s
 Full matrix would need 0.63 GB in float64; chunked version held one 250,000-row chunk at a time.
 Explained variance ratio, first 5 PCs: [0.511 0.1348 0.0958 0.0806 0.06]
 Cumulative, first 5 PCs : [0.511 0.6458 0.7416 0.8222 0.8822]
 Eigenvalues sorted descending: True

8. Summary of measured results

All numbers below are pulled from the runs above (not typed in by hand). They depend on the machine, so re-run the notebook on your own computer or Colab to see your own timings.

```
i_mil = sizes.index(1_000_000)
print(f"REAL DATASET ({n:}, x {d})")
print(f"  naive: {real_times['naive (loops + eig)']*1e3:.2f} ms | vectorized+eig: {real_times['vectorized + eig']*1e3:.3f}
print(f"\n1,000,000 ROWS")
for name in METHODS:
    t = scaling[name][i_mil]
    print(f"  {name:<22s} {'skipped (too slow)' if np.isnan(t) else f'{t:.3f} s'}")
print(f"  speedup of optimized (eigh) over vectorized+eig: {scaling['vectorized + eig'][i_mil]/scaling['optimized (eigh)']}|
print(f"\nFEATURES: at d = {dims[-1]}, eig is {t_eig[-1]/t_eigh[-1]:.1f}x faster than eig")
print(f"\nMEMORY at 1M rows: vectorized {mem_vec:.0f} MB | optimized {mem_opt:.0f} MB | in-place {mem_inp:.0f} MB | chunked
print(f"\nSTRESS TEST: {n_big:,} rows processed in {elapsed:.1f} s using chunks")
```

```
REAL DATASET (1,134 x 17)
  naive: 32.54 ms | vectorized+eig: 0.409 ms | optimized (eigh): 0.300 ms
```

```
1,000,000 ROWS
  naive (loops + eig)      skipped (too slow)
  vectorized + eig        0.213 s
  optimized (eigh)        0.313 s
  optimized float32       0.282 s
  svm                     1.720 s
```

Please ensure this tracking sheet is completed as a group, as it will be used to monitor your individual participation within the team.						
Team Members	Role/ Responsibility	Email				
Ariane Itetero	Performance and project packaging lead	a.itetero@alustudent.com				
Olga Ikirezi	PCA and mathematics lead	o.ikirezi1@alustudent.com				
Task Allocation and Tracker						
	Task Allocated	Assigned Member	Deadline	Completion Status (Not Started/In Progress/Completed/Did not participate)	Reviewed by Team? (Yes/No)	Comments
	Find/choose African dataset; data cleaning and preprocessing	Olga Ikirezi	29 Sep 2026	Completed	Yes	Updated 30 Sep
	Covariance matrix, eigenvalues and eigenvectors	Olga Ikirezi	29 Sep 2026	Completed	Yes	
	PCA projection, explained variance, dynamic selection of components, plots/results	Olga Ikirezi	30 Sep 2026	Completed	Yes	
	Task 3 optimization and benchmarking	Ariane Itetero	29 Sep 2026	Completed	Yes	
	Large-dataset testing	Ariane Itetero	29 Sep 2026	Completed	Yes	
	Notebook organization/testing	Ariane Itetero	30 Sep 2026	Completed	Yes	
	README and GitHub repository organization	Ariane Itetero	30 Sep 2026	Completed	Yes	
	Decide final number of PCs; explain trade-off and information lost for our African dataset	Olga Ikirezi & Ariane Itetero	30 Sep 2026	Completed	Yes	Depends on actual PCA results
	Final notebook review; make sure every code cell has visible output	Olga Ikirezi & Ariane Itetero	30 Sep 2026	Completed	Yes	
	Fill in contribution sheet; export the PDF	Olga Ikirezi & Ariane Itetero	30 Sep 2026	Completed	Yes	
Meeting Attendance Log						
Meeting Date	Agenda	Facilitator	Attendees	Key Discussion Points	Next Steps	
29 Sep 2026	Task distribution	Olga	Ariane Itetero, Olga Ikirezi	Divided the work: Olga owns PCA and the math; Ariane owns performance (Task 3, benchmarking, large-dataset testing) and packaging (README, GitHub, notebook). Agreed that the number of PCs, the trade-off, and the information lost are decided together.	Olga starts dataset selection and cleaning; Ariane starts Task 3 optimization and repo setup	
30 Sep 2026	Progress tracking	Olga	Ariane Itetero, Olga Ikirezi	Reviewed the task split and checked progress on each part.	Review both parts together, agree on the number of PCs and information-loss trade-off, and finalize the submission.	